

Benjamin Schindler

AI Engineering Lead & AI Engineer

Agentic AI Systems — RAG — Product AI — M.Sc. Data Science

+56 9 5627 9434 benjamin.schindler@gmail.com Santiago, Chile

benjaminschindler.vercel.app github.com/BenjaSchindler LinkedIn

Profile

AI Engineering Lead who builds and ships production LLM systems — multi-agent workflows, RAG, and the deterministic-agentic hybrid patterns that keep them safe in regulated domains. Owns systems end-to-end, from problem framing and evals to deployment and monitoring. Currently CTO at Doctor911 (early-stage healthtech), leading a 3-person product team; the AI products I designed and shipped contributed to a 10× increase in web traffic and the company reaching profitability.

Experience

- **Doctor911** — *early-stage healthtech* Mar. 2025 – Present
Chief Technology Officer & AI Engineering Lead (promoted) Jan. 2026 – Present
 - Lead a 3-person product team (2 engineers, 1 UX/UI designer); own the IT/AI infrastructure, the data platform, and the AI roadmap and technical standards (evals, observability, HITL safety, CI/CD).
 - Shipped the AI products underpinning the company's 10× traffic increase and path to profitability; promoted IC → CTO in under a year.
- **AI Engineer** Mar. 2025 – Jan. 2026
 - Built and shipped the company's conversational AI platform (LangGraph + FastAPI on GCP Cloud Run): a 4-agent WhatsApp system covering onboarding, sales, support, and order fulfillment (Meta Flows, WhatsApp catalogs, Transbank payments), and a 2-agent web system with a ~100 ms semantic intent router.
 - Built a form-based RAG recommender and an XGBoost medical exam recommender, exposed as shared tools across all agents.
 - Architected a deterministic-agentic hybrid pattern for the regulated side of the product: LangGraph orchestration constrained by structured business rules for payments, booking, and doctor-in-the-loop clinical approvals, with PII guardrails.
- **WiseConn Latam** Sep. 2024 – Mar. 2025
Data Science Intern (final-year M.Sc. internship)
 - Built CNN-based time-series forecasting for irrigation scheduling and classification models for solar installation analysis; deployed to production on AWS (SageMaker, S3) with validation and metric tracking, where they anticipate risk and reduce operating costs.
 - Led a cross-department data governance initiative consolidating the company's data.
- **Unitti** Jul. 2022 – Jul. 2024
Jr. AI Engineer
 - Built a Natural Language → SQL application (Flask, LLMs, Azure, GCP), improving SQL generation quality through user-feedback loops.
 - Developed Python APIs on PostgreSQL with automated testing, improving backend stability and query performance.

Selected Projects

- **MiAutoCheck** (client) — multimodal, multi-agent used-vehicle valuation API: a VLM inspection agent plus 5 parallel research agents (pricing, reliability, recalls, ownership cost, competitors) consolidated by a supervisor into PDF valuation reports. GPT-5-mini with Anthropic prompt caching to contain latency and cost; Tavily search.
- **EPE** (client) — workplace well-being platform with LLM safety guardrails (crisis / PII), per-user prompt A/B versioning, LLM-as-judge evals with drift monitoring, and Langfuse observability.
- **Doctor911 Internal Ops Agent** — built on the open-source OpenClaw framework to automate team coordination and recurring operational tasks.

Education

- **Universidad Adolfo Ibáñez** — *Master of Science in Data Science*, GPA 3.6 / 4.0 2024 – 2026
Thesis defended with Maximum Distinction. Thesis (*SMOTE-LLM*): hybrid LLM + embedding-space oversampling for imbalanced text classification; +2.25 pp macro-F1 over SMOTE (3,675-configuration benchmark, $p < 0.0001$).
- **Universidad Adolfo Ibáñez** — *Software Engineering* 2020 – 2025
- **Universidad Adolfo Ibáñez** — *Industrial Engineering* 2020 – 2025

Technical Stack

- **AI / ML:** LangGraph, LangChain, multi-agent orchestration, agent harness engineering (tool design, context management, guardrails), RAG, LLM evals (LLM-as-judge, regression suites), structured outputs, human-in-the-loop workflows. LLM providers: OpenAI, Anthropic, Gemini. Classical ML: XGBoost, PyTorch.
- **Platform & Cloud:** Python, TypeScript, SQL, FastAPI, Next.js, Docker, CI/CD; GCP (Cloud Run, Vertex AI, Cloud SQL), AWS (SageMaker, S3), Supabase, PostgreSQL; Langfuse observability.
- **Spoken Languages:** Spanish (Native), English (Advanced), German (Basic).