

Benjamin Schindler

AI Engineer & AI Engineering Lead

Agentic AI Systems — RAG — Product AI — M.Sc. Data Science

+ (56 9) 56279434 benjamin.schindlerv@gmail.com Santiago, Chile

LinkedIn Portfolio GitHub

Profile

AI Engineer who builds and ships production LLM systems — multi-agent workflows, RAG, recommendation flows, and the deterministic-agentic hybrid patterns that keep them safe in regulated domains. Owns systems end-to-end, from problem framing and evals to deployment and monitoring. Most recently CTO at Doctor911, where the AI products I designed and shipped contributed to a 10× increase in web traffic and the company reaching profitability.

Experience

- **Doctor911** Mar. 2025 – Present
Chief Technology Officer & AI Engineering Lead (promoted) Jan. 2026 – Present
 - Lead AI engineering and product across web, WhatsApp, and internal operations; own the AI roadmap and technical standards (evals, observability, HITL safety, CI/CD).
 - Shipped the AI products underpinning the company's 10× traffic increase and path to profitability; promoted IC → CTO in under a year.*AI Engineer* Mar. 2025 – Jan. 2026
 - Shipped two production multi-agent AI systems on LangGraph + GCP (Cloud Run, Vertex AI, CI/CD, monitoring; see *Selected Projects*), plus an XGBoost medical exam recommender wired as a shared tool across agents.
 - Architected a deterministic-agentic hybrid pattern combining LangGraph orchestration with structured business rules for payments, booking, and doctor-in-the-loop clinical decisions.
- **WiseConn Latam** Sep. 2024 – Mar. 2025
Data Science Intern
 - Built CNN-based time-series forecasting for irrigation scheduling and classification models for solar installation analysis (see *Selected Projects*).
 - Led a data governance initiative to organize company data across departments.
 - Integrated reproducible model workflows into AWS (SageMaker, S3) with validation and metric tracking.
- **Unitti** Jul. 2022 – Jul. 2024
Jr. AI Engineer
 - Developed Python APIs on PostgreSQL with automated testing, improving backend stability and query performance.
 - Built a Natural Language → SQL application (Flask, LLMs, Azure, GCP), improving SQL generation quality through user-feedback loops.

Selected Projects

- **Doctor911 Conversational AI Platform** — production multi-agent systems on LangGraph + FastAPI + Cloud Run spanning WhatsApp (4 agents: *welcome, commerce, support, async-order*) and Web (2 agents with a ~100 ms semantic intent router). Form-based RAG recommender and an XGBoost exam recommender wired as shared tools; integrates Meta Flows, WhatsApp catalogs, Transbank, Supabase, and Cloud SQL.
- **Direct Client Work** — two end-to-end AI products delivered for external clients. **MiAutoCheck:** multimodal, multi-agent used-vehicle valuation API (VLM inspection + 5 parallel research agents + supervisor consolidation → PDF reports; GPT-5-mini, Anthropic prompt caching, Tavily). **EPE:** workplace well-being platform with LLM safety guardrails (crisis / PII), per-user prompt A/B versioning, LLM-as-judge evals + drift monitoring, and Langfuse observability.
- **Doctor911 Internal Ops Agent (OpenClaw)** — internal agent built on the OpenClaw open-source framework to automate team coordination, recurring operational tasks, and management workflows.

Education

- **Universidad Adolfo Ibáñez** — *Master of Science in Data Science* GPA 3.6 / 4.0
Thesis defended with Maximum Distinction. Thesis (*SMOTE-LLM*): hybrid LLM + embedding-space oversampling for imbalanced text classification; +2.25 pp over SMOTE.
- **Universidad Adolfo Ibáñez** — *Software Engineering* GPA 3.5 / 4.0
- **Universidad Adolfo Ibáñez** — *Industrial Engineering* GPA 3.4 / 4.0

Technical Stack

- **AI / ML:** LangGraph, LangChain, multi-agent orchestration, RAG (embeddings + vector retrieval), tool-calling, structured outputs, human-in-the-loop workflows, recommendation flows, Natural Language → SQL, forecasting. LLM providers: OpenAI, Anthropic, Gemini. Classical ML: XGBoost, PyTorch, TensorFlow, Pandas.
- **Serving, MLOps & Cloud:** FastAPI, Docker, REST APIs, CI/CD, model & LLM evals, observability and monitoring, structured logging; GCP (Cloud Run, Vertex AI, Cloud SQL), AWS (SageMaker, S3), Supabase, PostgreSQL, Linux.
- **Languages & Frameworks:** Python, TypeScript, SQL/NoSQL, Next.js, React, Flask.
- **Spoken Languages:** Spanish (Native), English (Advanced), German (Basic).